Recipe Semantics for Noncognitivism

Abstract. Noncognitivists must account for the menagerie of moral attitudes of which we are psychologically capable. This paper offers a systematic explanation of these attitudes by means of a recipe semantics. Unlike extant noncognitivist theories, this proposal does not aim to justify the behavior of moral attitudes in terms of any underlying cognitive function. The recipe semantics allows us to characterize them in terms of the distinctive functions that their constituent components play in other contexts, while admitting that they may have no particular function in many. In order to make this account more palatable, the paper concludes by offering an account of the evolution of normative attitudes that would render their occasional non-functionality unsurprising.

Normative noncognitivism is a negative theory: it holds that normative judgments are *not* beliefs. Turning this into a satisfying positive metaethical theory requires saying much more about what moral attitudes actually are. To do so, noncognitivists have traditionally looked to the motivational capacity of moral judgments. They have focused on providing an alternative characterization of straightforwardly predicative moral judgments such as *insurance fraud is wrong*, *charitable donations are morally exemplary*, and *selfishness is sometimes permissible*. Such judgments are said to not really be beliefs, despite their surface form, because they do not represent the world as being any particular way. Instead, they reflect something about how we are moved to act.[1]

1A.J. Ayer (1936) and Simon Blackburn (1984) treat moral judgments as a special sort of approval and disapproval. Charles Stevenson (1937) described them as a variety of interest or partiality. Allan Gibbard (2003) compared them to plans. Mark Timmons and Terry Horgan (2006) suggested that they are a special non-representational species of belief. Mark Schroeder (2008) proposed (though refrained from himself endorsing) that they should be thought of as attitudes of favoring or disfavoring blame for certain actions. These views all place great weight on the pressures moral judgments put on us to act in certain ways. Some of the

However, not all moral attitudes are straightforwardly predicative moral judgments. We not only judge that insurance fraud is wrong and charitable donations are exemplary, we also can judge that providing misleading responses to inappropriate questions is not wrong and that eating meat is morally permissible if animals are not conscious.

Questions over the nature of logically complex moral judgments have been much discussed as part of the Frege-Geach problem, but logical complexity is just one sort of complication found in moral attitudes. An adequate positive theory would also include a characterization of a range of other attitudes.

In this paper, I will sketch a strategy for characterizing a variety of moral attitudes. My approach will differ from extant noncognitivist views in that I don't aim to rationally vindicate the distinctive features of other moral attitudes. Instead, I aim only to provide an account of what the attitudes are and how they might have come to be. I will start by cataloging some of the complexities that must be dealt with. Next, I will present an account of our propositional attitudes that will play an important role in shaping how I propose we characterize moral attitudes. Third, I will offer what I call a "recipe semantics" for characterizing the other moral attitudes. This approach seeks to characterize the other moral attitudes in terms of the distinctive functions that their constituent components play in other contexts. I will conclude with an account of the cognitive evolution of normative judgments that makes a recipe semantics especially reasonable on the grounds that many moral attitudes may lack a specific function.

1. A Menagerie of Moral Attitudes

---

remaining difficulties involved in providing an adequate characterization are explored in David Merli (2008).

Cognitivists hold that judgments such as (A) can be characterized by their content.

A) It is wrong to collect trophies of endangered species.

The judgment expressed by (A) takes a proposition about the distribution of a certain property (*wrongness*) as its object. Cognitivists disagree about what this property is, but they agree that there are some definite conditions that must be met for (A) to be true.

Noncognitivists, in contrast, hold that such judgments have no special propositional content. There is no property of wrongness (at least not in any robustly metaphysical sense) that we ascribe to acts in judging them wrong. Rather, in judging an act to be wrong, noncognitivists typically think that we adopt a certain conative attitude towards it. What makes these attitudes moral judgments is not their representational content, but their motivational force.

An account of straightforwardly predicative moral judgments won't tell us much about the nature of moral judgments that involve logical complexity. Consider (B), which employs 'wrong' under the scope of a negation.

B) It is not wrong to collect trophies of extinct species.

(B) doesn't express an attitude that is motivational in quite the same way as (A), so an account of predicative moral judgments will not carry directly over to negational judgments. We are not moved to do anything in particular solely because we believe that something is *not* wrong.

This problem with negation is only the most basic form of logical complication.[2] The same problem also arises for conjunction and

2 Stemming from pivotal work by Peter Geach (1965) and John Searle (1962), the question of logical complexity has taken a prominent place in discussions of noncognitivism. Though it is typically fit under the inclusive label of *The* Frege-

disjunction, as in the attitude expressed by (C), and quantification as in the attitudes expressed by (D) and (E).

> C) Human beings either have a special moral status or we should not be testing on animals.
>
> D) Everything you did today was worse than the thing you did before it.
>
> E) Most killings are deeply morally wrong.

Quantification hasn't received the same attention as conjunction, disjunction, and negation, but it is equally problematic for noncognitivists. We can make quantitative moral judgments about an issue without making any straightforward predicative judgments about it.

In addition to ordinary quantification, moral judgments can also take plural and generic forms, such as are expressed by (F) and (G). Generics are especially interesting because of their unsystematic behavior. Sarah-Jane Leslie (2007) has argued that this owes to the fact that generics employ a cognitively basic form of generalization, which is sensitive to the functional needs of such generalizations, and doesn't admit of clean truth conditions.

> F) The wrongs that the colonists did to the natives were more varied and numerous than the wrongs that the natives did to the colonists.

Geach Problem, that problem encompasses a variety of different issues. One part of the problem – the part that I focus on in this paper – is to explain what the logical complexity of moral judgments amounts to. The other parts of the problem involve explaining why logically complex attitudes relate to each other in the ways that they do, how such as how attitudes can be inconsistent, and how they can figure into arguments. The challenges involved in these problems are explored in depth in (Schroeder 2008). I have discussed them in *Author (xx 1)*, and the solution I present there for this second part of the Frege-Geach problem is compatible with the solution I present here for the first part of the Frege-Geach problem.

G) Wrongs are a stain upon the moral character of decent people.

Geach, along with many subsequent commentators, was especially interested in conditionals, such the attitude expressed by (H), that play an important role in moral arguments.

H) If lying is wrong, then getting your little brother to lie is wrong.

These conditionals can be given a truth-functional interpretation in which they are true unless the antecedent is true and the consequent is false. However, this interpretation is widely thought to be inadequate for a number of indicative conditionals. The indicative conditionals found in expressions of moral judgments like (H) are subject to all the same considerations that motivate non-truth conditional interpretations of indicative conditionals in other contexts. If such an interpretation is correct, then they present a different metaethical problem from negation, disjunction, conjunction, and quantification.

In addition to purely logical complications, moral judgments can take the form of subjunctive conditionals, as expressed by (I), and can involve tense, as expressed by (J).

I) If it had been wrong to bring children into this world, we would not have done it.

J) Abortion is as wrong today as it was fifty years ago.

Noncognitivists must also account for the existence of mixed normative judgments. Mixed normative judgments, as expressed by (K) and (L), combine different flavors of normativity into a single attitude

K) You ought (epistemically) to know that what you did was wrong.

L) You should (rationally) keep track of what you shouldn't (morally) do.

Finally, I (xx 2) have argued that noncognitivists owe an explanation of moral attitudes other than judgment. Not only can we judge that actions are wrong, but we can desire that they be (or not be) wrong. We can regret that we acted immorally, and hope that we chose the right thing to do. We can wonder whether utilitarianism or deontology is correct. We can find some moral principles intuitive and others unintuitive. We can imagine fictional scenarios where different moral principles hold. We can suppose novel or bizarre moral principles for the sake of arguments.

We must make sense of these other attitudes just as we must make sense of straightforwardly predicative moral judgments. An account of what it is to judge that an action is wrong will not tell us what it is to hope that or wonder whether we did the right thing.

The challenge facing noncognitivists is to provide a suitable account of all of these other moral attitudes. That account must make sense of negative, conjunctive, disjunctive, quantificational, conditional, tensed, and mixed normative moral judgments and also explain what is going on in the case of moral desires, regrets, hopes, uncertainty, intuitions, imaginings, and suppositions.

Ideally, the noncognitivists' account would offer a systematic characterization of all of these different attitudes. After laying out some assumptions about the cognitive structure of attitudes, I will return to give a sketch of one systematic account.

2. Context and Meaning Determination

Since moral attitudes have so much in common with other propositional attitudes, insight into former might come from the latter. In this section, I will present one view about cognitive structure of ordinary

propositional attitudes. This view will incorporate substantial and controversial assumptions, but it is doubtful that we can develop an adequate theory of moral attitudes without making a few assumptions about how our minds work.

On the view that I will present, our propositional attitudes involve relations to cognitive representations whose representational contents are determined by the representational properties of their parts. Those parts can be individuated in terms of non-semantic properties, and they receive their representational content from their designed function in certain contexts. I will describe each of these components piece by piece, and then sum up.

The Representational Theory of Mind

According to the Representational Theory of Mind (RTM)[3], ordinary propositional attitudes involve relations to mental representations whose content is the propositional object of the attitude. What sort of attitude it is (belief, desire, etc.) will generally depend on the relation taken to the mental representation. To believe a proposition, for instance, might be to be disposed to token representations of that proposition in a certain way and use those representations in deciding what to do with the goal of getting what one wants.[4]

In order for this theory to go beyond merely unpacking our concepts of propositional attitudes, "mental representation" must be

3 This term was coined by Jerry Fodor, who uses it to describe a constellation of theories (1998). I will restrict its usage to the core view as defined above. (A view that is found, for instance, in Field 1978.) Unlike Fodor, I will use Representational Theory of Mind to name a theory distinct from the Language of Thought Hypothesis, which makes additional assumptions about the structure of representations, and the Computational Theory of Mind, which makes additional assumptions about the way that representations are handled in deliberation and reasoning.

4This is intended to allow for the possibility that the relation is dispositional. Thus, in order to believe a proposition, we need not actually represent that proposition, so long as we are still suitably related to possible representations of it. See Schwitzgebel (2002) for a sophisticated account of this.

understood in a robust and realistic sense. We may discover that our minds do not employ anything that counts as a mental representation in this robust sense. It would still be true that beliefs involve representations in a pleonastic sense, but RTM would then be false.

Those who accept RTM owe an account of how mental representations can represent propositions. For simplicity, I will assume that a teleological theory is correct. Roughly put, mental representations get their representational contents by virtue of being designed, in some sense or other, to have the function of representing propositions.[5] This theory won't differ from other plausible theories of cognitive representationality where it counts for the our present purposes, and so readers may substitute their own preferred psychosemantic theory wherever appropriate.

Compositional Representations

Representations can be either structured or unstructured. Structured representations are built out of constituents that themselves have representational contents. I will refer to the ultimate constituents of individual mental representations as "concept tokens", and reserve the term "concept" for types of concept tokens that are individuated by their designed functions.[6]

Sentences in natural language are structured, since many of the individual phrases and the words that constitute them represent properties, entities, or propositions. Maps have a different kind of structure, insofar as the different sub-regions of the map themselves have representational content. A line might represent the coast and the way it is positioned within the whole map represents the relation between the coastline and other points of interest.

5Variants of this view are discussed by Fred Dretske (1995), Ruth Millikan (1984), and David Papineau (1984).

6This usage is inspired by Laurence and Margolis's (2007) proposal.

Structured representations are typically compositional, meaning that the representational contents of the parts help to determine the representational contents of whole representation. The representational contents of concepts tokens can't themselves determine how it is that concept tokens can neurophysiologically be combined into a whole, though they will have a say in whether the combinations they enter into are meaningful.[7] Non-semantic properties must determine how concepts can be strung together in the mind. These properties, whatever they are, are "syntactic". As a result, there may be some syntactically viable combinations of concept tokens that are semantically incoherent – viable combinations assembled without purpose.

Orthographic Identities

By supposition, the representational content of a mental representation depends on its designed function. The designed function of a representation may result from the designed function of its constituent concepts. A representation may be designed for a particular representational function by virtue of the fact that its constituent concepts were designed to be used to create representations. This resembles the situation in natural language semantics, where the meaning of a sentence is determined by the meanings of its words, and the meanings of words are in turn determined by the kinds of contributions they make to the meanings of sentences.

By supposition, concept tokens get meanings (i.e. get to be tokens of a given concept) from their designed functions. Designed functions depend upon behavior in some range of contexts, and so it is necessary that concept tokens have some identity across contexts that does not depend upon their representational content. Consider how difficult it would be to assign contents to concepts without being able to recognize them across contexts. Quine's radical interpreters at least had

7See Stich (1983).

the identity of the sound ‘gavagai’ to work with. If they hadn’t – if the word had sounded arbitrarily different in each context of use – their task (and language itself) would have been utterly impossible. So it is with cognitive semantics. The difficulty is not merely epistemic; a concept token must have a non-semantic identity in order to support a robust enough pattern of uses to count as having a designed function under any reasonable psychosemantics.

Since it is used to assign representational contents, the property whereby a concept token is identified from context to context must be non-representational. Given the analogy with written language, I will refer to such a property, whatever it is, as “orthographic”. Orthographic properties might be closely associated with syntactic properties (the properties that determine how concept tokens can be combined with others), but they need not be the same.

Spandrel Contexts

Our concepts may or may not have a designed function in every context in which their syntactic properties allow them to occur. In biology, phenotypical trait occurrence often outstrips designed function. Natural selection exerts imprecise forces; adaptations that are helpful in one context may have effects elsewhere. Our bodies have a variety vestigial and by-product traits. We should expect this to be as true for cognition as it is physiology.[8] I will refer to contexts in which a concept type occurs in which it has no designed function as a “spandrel context”.

Suppose that that we identify the concept *horse* as a concept whose designed function is to help track the properties of horses by figuring into representations of horses. It isn’t part of the designed function of the concept *horse* to help distract us when we are bored, but

8In fact, given the extremely short time in which human cognition has developed from our vastly cognitively inferior forebears, it would be unsurprising if it is especially messy and contain many states that serve no important or meaningful function as a byproduct to other important advances.

its tokens can be employed for this purpose. This use need not go against the concept's designed function -- the concept may not have been designed *not* to be used as a distraction. Its function may be given by what it is supposed to contribute in the context of certain kinds of attitudes (beliefs and desires). It may have no functions in others, such as daydreams. If this is the case, then these other contexts are spandrel contextes for the concept. Its token's syntactic properties allow them to appear in those contexts, but it isn't designed for those contexts.

3. Recipe Semantics

In the last section, I put forward a view on which propositional attitudes involve mental representations. I suggested that individual mental representations can be identified across contexts by virtue of their orthographic properties just as we identify written words in different contexts by their shape and spoken words by their sound. My proposal for handling the other moral attitudes will depend on extending this approach to moral attitudes. I will start by generalizing the idea of mental representations into a class of mental entities that need not have representational properties. Then I will provide a recipe semantics for daydreams to serve as an example for the view I hold about moral attitudes. Finally, I will combine the two ideas to provide a recipe semantics of non-straightforwardly predicational moral attitudes.

Presentations

The orthographic properties of non-mental representations can be divided up into resemblance classes. Written words, for instance, are individuated by their shapes. They look distinct from each other, but they share more in common with each other than they do with spoken words. An *orthographic class* is a set of entities whose orthographic properties resemble each other and which includes a system of

representations as a subclass. There is an orthographic class for sentences.[9] There are others for maps, tables, and graphs.

Members of an orthographic classes are members of that class solely by virtue of the similarities among the non-semantic properties that they share with other members. There is no reason to think that each member of an orthographic class must be representational. It is easy to find non-representational members of orthographic classes of non-mental representations. There are sequences of letters, sounds, and collections of squiggles on a two dimensional field, that have no representational significance. Take for instance, the final stanza of Lewis Caroll's Jabberwocky:

'Twas brillig, and the slithy toves

Did gyre and gimble in the wabe:

All mimsy were the borogoves,

And the mome raths outgrabe.

This belongs to an orthographic class that also contains natural language sentences in English, but since the words included have no established meaning, they do not represent anything.

I will use the term "presentation" to refer to the members of an orthographic class. Representations are presentations, but not all presentations are representations. The stanza from Jabberwocky consists in presentations which do not rise to the level of representations. Since there are kinds of meaning that are not representational, there may be some meaningful natural language

---

9 Orthographic classes may be highly vague. There needn't be one single precise orthographic class for sentences. Thus, there needn't be a clear cut answer as to whether or not sentences in English and Chinese fall into the same orthographic class.

presentations that are not representational and have no representational significance.[10]

If the mind utilizes representations in having propositional attitudes, it may also utilize non-representational presentations in other attitudes. Mental presentations are conceptually coherent. They are simply presentations that fail to meet a contingently-satisfied condition for counting as a representation. On the present assumptions, any presentation which was not designed to function as a representation will count as non-representational.

The Presentational Theory of Mind

According to RTM, propositional attitudes are relations to mental representations. If beliefs involve a relation to mental representations, then it is not far-fetched that moral judgments should involve relations to mental presentations. Beliefs and moral judgments are very similar (Horgan and Timmons 2006). It would be surprising that they should work fundamentally differently at the cognitive level.

If beliefs and moral judgments involve relations to presentations, then we can locate the difference between beliefs and moral judgments either in the kinds of presentations they employ, the relations they involve to those presentations, or both. I will take it for granted that the primary difference between ordinary beliefs and ordinary desires is the former. To believe something and to desire it involve taking different relations to the same representation.

Now to my proposal: *straightforwardly predicative moral judgments are attitudes that involve relations to complex structured presentations that contain non-representational concepts whose proper function is not to represent but instead to influence action*. These

10 There are many words, including logical connectives, that don't represent anything themselves but still have representational significance.

concepts are moral concepts, and they are applied in just the same way as their representational analogues. Since the designed function of neither the whole presentation nor its constituent moral concepts is to represent, neither is representational. But since the designed function conveyed to presentations by the presence of the moral concept is to direct action, they are also not without a designed function altogether.

We can go on to characterize the designed function of moral concepts in a variety of different ways, so this proposal does not tell us exactly what sorts of attitudes straightforward moral predicates are. We haven't said anything about the important question of how to distinguish moral judgments from other normative judgments, or indeed from other motivational attitudes. However, extant proposals can be carried over rather straightforwardly to the present context.[11]

A Recipe Theory of Daydreams

Philosophers of mind have traditionally focused on attitudes, such as belief and desire, that play an important role in influencing our behavior. Not all of our attitudes are like this. Daydreams, for instance, are not. This means that many strategies applied to characterizing attitudes cannot always be carried over. We cannot make sense of what it is to daydream about a certain situation in terms of what that daydream does in our mental life, as functionalists would ordinarily hold. In order to understand these states, we may need to approach them as byproducts of other cognitive faculties.

11 It might be objected that my proposal will restrict moral judgments only to those creatures who employ presentations in cognizing. Intuitively, insofar as it is epistemically possible that we do not, we should not make it necessary for a creature to do so in order to make moral judgments. I am sensitive to this worry, and I will allow that moral attitudes may not always require presentations. Nevertheless, I don't think it would be misguided to limit our focus onto to creatures like us. One way to have moral attitudes – maybe not the only way – is to have attitudes to presentations that include moral concepts.

I propose that daydreams, and in particular their representational contents, ought to be understood in terms of their parts, and that those parts in turn are to be characterized in terms of how they are used elsewhere. We can't recognize a concept token as representing a horse, for instance, solely in terms of its deployment in daydreams. Instead, we must look to how the same concept token is (supposed to be) used in other cognitive contexts.

The existence of orthographic properties and spandrel contexts are important when it comes to understanding the nature of daydreams. A dream gets its content by employing concept tokens that play an important role in other contexts. These concept tokens are recognized by their orthographic properties. A concept token of a horse, as it occurs in a dream, represents a horse only insofar as it was designed to have a certain representational content in other contexts. The fact that we are dreaming of a horse, and not a teapot or Mars, depends on the fact that the same orthographic representation has the function of keeping track of the properties of horses on other occasions.

A Recipe Semantics for Moral Attitudes

In a previous section, I suggested that straightforwardly predicative moral judgments involve a special kind of mental presentation. In the present, I will explain how to extend this characterization to handle more complex moral attitudes. I will start by giving an overview of the strategy, and then go on to demonstrate its application in several examples.

Moral concepts are characterized by a non-representational proper function that operates in the context of straightforwardly predicative moral judgments. In order for a concept token to be an instance of a moral concept, it must be designed to be used in a certain way in certain contexts. Given that a concept token can be recognized

in a variety of contexts based on its orthographic properties, and given that it need not have a designed function in all of the contexts in which it occurs, it is easy to systematically extend an account to handle each member of the menagerie of other attitudes.

The general strategy is to characterize attitudes in terms of a recipe for constructing them (a "recipe semantics"). The recipe will describe the ingredients that must be combined to produce the attitude. The ingredients may be characterized by their designed function in other contexts. There is nothing special about these attitudes, other than the parts from which they are made. They may lack a distinctive function or important role in cognition.

By following this approach, we avoid the need to describe the function of these other attitudes. This makes this approach very different from traditional approaches to the problem. Unlike Simon Blackburn and Allan Gibbard, I do not think we need to vindicate the forms our attitudes take. We do not need to explain why it is that they behave as they do in terms of some essential function. Some of our attitudes may not play any important role in our psychologies. Not every aspect of human psychology needs to have an intelligible rationale.

On my proposal, complex moral attitudes are like daydreams. Dreams get their particular contents because they involve representations whose characteristic function is restricted to other contexts. Complex moral judgments will often involve concepts whose function is specific to other contexts.

Negation

In order to see how this proposal works in practice, I will start by applying it to one of the simplest case: negation. Negational moral judgments are the sorts of judgments that we typically express with negations of moral predicates. They have the appearance of logical

inconsistency with straightforwardly predicative moral judgments. The judgment that eugenics is not wrong is an example of a negational moral judgment. It appears to be inconsistent with the judgment that eugenics is wrong.

It has been common to characterize negational attitudes by their functional roles; they are states that are *inconsistent*, in some way or other, with straightforwardly predicative moral judgments (Blackburn 1988). And, according to the traditional story, the incompatibility arises from the attitude's functional role. For instance, the functional role of both straightforwardly predicative and negational moral judgments might be to direct us to make a certain state of affairs a reality. If incompatible attitudes are directed to realizing incompatible states of affairs, then they cannot be jointly satisfied, and might therefore be deemed inconsistent.

I am not pretheoretically opposed to ascribing straightforward predicative moral judgments and their negations inconsistent functional roles, but I do not believe that noncognitivists should feel obliged to do so. Instead, noncognitivists can simply state that negational moral attitudes involve presentations with a special sort of negation concept.

In order to make sense of the appearance of logical inconsistency among representations, it is helpful to attribute them some structure (Baker and Woods 2015). This can be done through the presence of a distinctive concept in the cognitive presentation. A negation concept is a concept which has a certain function in representations. That doesn't mean it can't occur elsewhere. *Horse* concepts plausibly function to keep track of horses and their properties, but they also occur in daydreams. The negation concept, I propose, is similar.

The characteristic function of a negation concept may be to flip the propositional content of cognitive representations, so that a

representation governed by a negation concept represents a proposition with the opposite truth value of the proposition represented by its negated content. This designed function of the concept is specific to representations. The concept cannot have this function in all contexts. In the context of a moral judgment, the negation concept has no specific function. It simply adds another bit to the presentation. In those other contexts, the existence of the negation concept may influence the behavior of the attitude in our cognitive lives. This is what produces the appearance of inconsistency. But its influence isn't a matter of that concept's function in that context.

A negational moral judgment involves a presentation that includes a negation concept as a constituent. Though the negation concept is designed to play a certain function in the case of representations, it need not play that specific function in the case of other presentations and so it may have no designed function whatsoever in the context of moral presentations. What identifies the moral judgment as a negational moral judgment is simply the presence of a constituent which has a certain function in other contexts.

Generics

Generics pose a difficult problem for linguists. On the one hand, generics seem to behave like ordinary quantifiers such as 'all' or 'most', on the other hand, they defy straightforward analyses in terms of quantities.

One of the most promising approaches to understanding generics treats them as a human cognitive default form of generalization.[12] The proper applications of generics are guided by purposes of this default generalization, which serves to keep track of important properties of a kind without requiring any specific frequency among that kind. Where it

12See Leslie (2008).

is advantageous to associate the kind with that property, the generic is properly applied.

We may assume that generics are used express mental states with the default mode of generalization, and that such modes are orthographically distinctive. Perhaps we employ particular generic quantificational concept tokens. But for something as basic as generic generalizations, the orthographic property may come from the structure of the presentation rather than from a proper constituent of it. Either way, there is something non-semantic that is distinctive about the form that generic judgments take. This distinctive property can be identified with generic quantification by virtue of designed function in certain representations.

The existence of moral generics can be explained in turn by the fact that the same distinctive orthographic traits that underlie generic judgments can be found among presentations in combination with moral concepts. Generic concepts or generic structures that are combined with representational concepts to have a certain representational role might also be combined with non-representational concepts. In such cases, they need not have any distinctive designed function. They can instead by their distinctive function in making default generalizations in other contexts.

Hope

According to RTM, propositional attitudes involve relations to mental representations. Typically, differences in kinds of propositional attitudes correspond to differences in the involved relationships to mental representations. To believe something and to desire it involve taking different relations to representations with the same representation. Just as for concept tokens, we can expect that these relations have a non-semantic identity. There must be some cognitive or neurological

features that underlie the behavior of the attitude and help us to recognize the relation across contexts.

If hopes are a distinct variety of propositional attitude and not a mere amalgamation of other more basic propositional attitudes, then we can expect that there will be a separate orthographically distinctive property that characterizes the relations we have to the representations we hope to be true. Even if hoping is characterized in terms of its distinctive functional role within our cognitive lives, the relation that plays that role is likely to be orthographically distinctive. Indeed, to have a robust functional role it must be recognizable across contexts, and so it must have some orthographic identity.

If our minds employ moral presentations and we are related to them in the same orthographically distinctive way that we are related to the representations that we hope, then we might say that we hope them as well, even though they lack representational content. Thus, the hope that lying was permissible could simply be interpreted as the attitude that results from adopting the relation characteristic of hopes to a presentation with a certain moral concept. Again, we supplement a recipe for creating the attitude with a characterization of the attitude itself. To create a moral hope, you must combine a presentation that includes a moral concept with the orthographic properties special to hope.

4. Belief, Exaptation, and Moral Judgments

The value of this recipe semantics depends on whether moral concepts are able to appear in many contexts in which they have no designed function.[13] There is nothing incoherent about moral concepts appearing where they lack a function, but it may still be thought that it is better to

13 In such cases, it is tempting to say that they lack a meaning. However, this does not follow. They certainly do not have a reference. And they do not make a specific contribution to the designed function of the attitudes in those contexts, but they are still meaningful by virtue of their contributions in other contexts.

avoid postulating mental states without designed functions. After all, there is some mystery as to why we have mental states that play no important role in our cognitive lives.

The reason that we are able to combine concept tokens and orthographically individuated relations in the way that we do depends ultimately on their syntactic properties. The syntactic properties of a concept token are those properties, whatever they are, that determine and how those concept tokens may be deployed in presentations. The semantic properties, which depend partly on the historical or counterfactual behavior of the state, cannot themselves explain why it is that we are able to form those presentations that we can.

There is no special burden that noncognitivists must bear in explaining the attitudes we're capable of having. Everyone must provide some explanation for why we have the attitudes that we do. Cognitivists face the challenge of explaining how we are capable of forming the diverse beliefs that we do. Jerry Fodor has passionately argued for a solution that looks to the structure rather than the content of our attitudes for an explanation (Fodor and Pylyshyn 1988, Fodor 1998, 2008). Fodor explained how we have the attitudes that we do in terms of how they can be built out of recombinable units. This explanation relies on syntax, rather than semantics. There is little reason to think the challenge of accounting for our diverse attitudes (or its solution) will look much different noncognitivists than for cognitivists.

Nevertheless, it might be doubted that our moral concepts have syntactic properties that allow them to be combined into attitudes which have no designed function. Why is it that we're able to hope that what we did was right, if the concept *right* doesn't make any specific contribution to the attitude?

The evolutionary history of moral attitudes might render these syntactic properties unsurprising. What follows is one story of how this history may have gone. It isn't the only viable story, but it goes a long way to warding off serious worries that arise from syntactic properties of moral concepts.

The key idea behind the story of the evolution of moral attitudes is that moral judgments evolved to help us to coordinate on social expectations.[14] By discussing norms, we learn how to avoid transgressing them. Such a system of coordinating attitudes must not have sprang up over-night. Chimpanzees, our closest relatives, have nothing clearly recognizable as judgments about morality.[15] Therefore, our ancestors must have started making such judgments only in the last several million years. Given that our species seems to have acquired much of its capacity for abstract reasoning, complex language, and cultural adaptation over the last 100,000 years, moral attitudes must have evolved quickly.

Our pre-moral ancestors would have felt pressures for coordinating social expectations. These pressures would have set to work on a cognitive system with a stock of non-moral attitudes including beliefs and desires. Moral judgments might have emerged sui generis from the cognitive void, but it seems more likely that they would have been spun off from either beliefs or desires.

On the present story, moral judgments resulted from alterations to beliefs. Our pre-moral ancestors might have started with ordinary beliefs about social rules and regulations. They could judge actions according to their social appropriateness within their culture. Generation

---

14 See (Gibbard 1993) and (Björnsson and McPherson 2014) for thoughtful presentations of this perspective.

15 They do seem to make judgments about fairness, and have expectations about how other chimpanzees will act.

by generation, they became increasingly disposed to form attitudes that functioned not to simply represent social standards, but to motivate them to live up to a set of standards that might differ slightly from those presently accepted by society, until the designed functions of these attitudes fell outside of the range of representationality.

This shift from representation to motivation could have resulted from several different pressures. The first of these is norm internalization. **Norm internalization** is the process by which individuals come to treat prevailing social norms as if they were reasons, and give them weight in deliberation independently of their social significance. Knowing and communicating social norms is very important for social animals insofar as it allows us to stay within our cultures' bounds. However, internalizing social norms and being transparently governed by them offers different benefits. Someone who has clearly internalized a society's regulations is more trustworthy and is less likely to receive serious punishment after norms are transgressed.[16]

A second force might have come in the form of a **progressive shift**. As culture flourished, social expectations must have become more flexible. With societies becoming more open to normative change, attitudes whose former function was tracking expectations might have acquired a new role in influencing them. Moral deviants had attitudes that were not defective, but progressive. Their attitudes might have played an important role in altering the existing norms. This non-representational function could have been established as the designed function insofar as it was selected for to help guide communities to framing and adopting new norms.

---

16 David Gauthier (1986) has given a vigorous defense of the theoretical value of norm internalization. Christopher Boehm (2010) describes some anthropological evidence that norm internalization is rewarded by leniency, when those norms are violated.

If our moral judgments resulted from exaptation of beliefs in line with norm internalization or a progressive shift, and moral concept tokens incorporated a motivational function in place of a representational function, then it would not be too surprising that they should have retained the syntactic properties common to beliefs and representational concepts. Our concept tokens of rules and social regulations combine normally with negation, conjunction, quantification, generics, conditionals, tense, and other attitudes*. If our moral concept tokens exapted from representational concepts, then they should continue to share some of their syntactic properties.*

This evolutionary story is far from the only story to be told about where our moral attitudes come from, and it may well be false. However, its plausibility undermines concerns about the syntactic properties that allow moral concepts to appear in the contexts they do. The fact that moral judgments do display the same syntactic properties as beliefs is not a major objection to the view.

5. Conclusion

In this paper, I have argued that noncognitivists can provide a recipe semantics for a variety of moral attitudes. Moral judgments are special insofar as they involve moral concepts. Moral concepts are characterized by the fact that they play a certain motivational role in the context of moral judgments. Moral concepts can be recognized in other contexts by their orthographic properties. The other attitudes, including both complex moral judgments and non-judgment attitudes, are characterized in terms of their orthographic similarities to other attitudes.

I conclude by saying something of the significance of this view for the practice of thinking and arguing about morality. It is not an error theory in the vein of J.L. Mackie (1977) and Richard Joyce (2001), though

it shares something of their spirit. Error theories hold moral judgments to be representational attitudes that represent the world in necessarily false ways. They suggest that our moral attitudes are radically mistaken. The recipe account that I have developed does suggest that many of our moral attitudes are in some sense accidental insofar as they were designed to play no particular function.

While they may be accidental byproducts of other functional states, nothing that I have said implies that they are *mistakes*. Whether or not it is a mistake to hold these attitudes is a normative issue. We can't directly infer from the absence of a special function that we should stop using them. We still use them in a variety of ways. Even if we can't ultimately find anything useful for them to do, it no more follows that we should stop making complex moral judgments, or give up our moral hopes, than that we should have our appendixes surgically removed. Ultimately, normative noncognitivists must decide whether these attitudes are mistaken by assessing how it is we feel about them.

Bibliography

Author (xx 1) [redacted]

Author (xx 2) [redacted]

Ayer, A. J. (1936). *Language, Truth and Logic*. London, V. Gollancz, Ltd.

Baker, Derek & Woods, Jack (2015). How Expressivists Can and Should Explain Inconsistency. *Ethics* 125 (2):391-424.

Blackburn, Simon (1984). *Spreading the Word*. Clarendon Press.

Blackburn, Simon (1988). Attitudes and Contents. *Ethics* 98 (3):501-517.

Björnsson, Gunnar & McPherson, Tristram (2014). Moral Attitudes for Non-Cognitivists: Solving the Specification Problem. *Mind* 123 (489):1-38.

Boehm, Christopher (2010). *Moral Origins: The Evolution of Virtue, Altruism, and Shame*. Basic Books.

Dretske, Fred (1995). *Naturalizing the Mind*. MIT Press.

Field, Hartry (1978). Mental Representation. *Erkenntnis* 13 (July):9-61.

Fodor, Jerry A. (1998). *Concepts: Where Cognitive Science Went Wrong*. Oxford University Press.

Fodor, Jerry A. (2008). *Lot 2: The Language of Thought Revisited*. Oxford University Press.

Fodor, Jerry A. & Pylyshyn, Zenon W. (1988). Connectionism and Cognitive Architecture. *Cognition* 28 (1-2):3-71.

Gauthier, David P. (1986). *Morals by Agreement*. Oxford University Press.

Geach, Peter T. (1965). Assertion. *Philosophical Review* 74 (4):449-465.

Gibbard, Allan (2003). *Thinking How to Live*. Harvard University Press.

Horgan, Terry & Timmons, Mark (2006). Cognitivist Expressivism. In Terry Horgan & Mark Timmons (eds.), *Metaethics After Moore*. Oxford University Press 255--298.

Joyce, Richard (2001). *The Myth of Morality*. Cambridge University Press.

Leslie, Sarah-Jane (2007). Generics and the Structure of the Mind. *Philosophical Perspectives* 21 (1):375–403.

Margolis, Eric & Laurence, Stephen (2007). The Ontology of Concepts: Abstract Objects or Mental Representations? *Noûs* 41 (4):561-593.

Mackie, J. L. (1977). *Ethics: Inventing Right and Wrong*. Penguin.

Merli, David (2008). Expressivism and the Limits of Moral Disagreement. *Journal of Ethics* 12 (1):25 - 55.

Millikan, Ruth G. (1984). *Language, Thought and Other Biological Categories*. MIT Press.

Papineau, David (1984). Representation and Explanation. *Philosophy of Science* 51 (December):550-72.

Schroeder, Mark Andrew (2008). *Being For: Evaluating the Semantic Program of Expressivism*. Oxford University Press.

Schwitzgebel, Eric (2002). A Phenomenal, Dispositional Account of Belief. *Noûs* 36 (2):249-275.

Searle, John R. (1962). Meaning and Speech Acts. *Philosophical Review* 71 (4):423-432.

Stevenson, Charles Leslie (1937). The Emotive Meaning of Ethical Terms. *Mind* 46 (181):14-31.

Stich, Stephen P. (1983). *From Folk Psychology to Cognitive Science*. MIT \Press.